

Preface

Welcome to the AAAI 2026 Workshop on *Bias in Multimodal AI: Representation, Risk, and Repair*.

As multimodal AI systems become increasingly central to how we interpret, generate, and act on information across text, image, audio, and other modalities, questions of bias have become more urgent and more complex. Bias in multimodal systems is not merely a matter of skewed datasets or isolated model failures. It emerges through interactions among representations, training objectives, interfaces, deployment settings, and social context. Understanding these interactions requires perspectives that cut across machine learning, natural language processing, computer vision, ethics, human-centered AI, and domain-specific applications.

This workshop was created to provide a forum for that conversation. Our aim is to bring together researchers and practitioners who study how bias is represented, amplified, diagnosed, and mitigated in multimodal AI systems. The theme of this year's workshop—*representation, risk, and repair*—reflects three complementary challenges: understanding how multimodal models encode socially consequential distinctions, identifying the technical and human risks that follow from those representations, and developing principled methods for analysis, intervention, and repair.

The 2026 program reflects the breadth and urgency of this topic. It includes keynote talks by Prof. Eduard Hovy and Prof. Sang-Wook Yi, an invited talk by Dr. Kyungreem Han, a panel discussion on representation, risk, and repair, and presentations of selected accepted papers addressing multimodal bias from technical, conceptual, and application-oriented perspectives. Together, these sessions highlight the importance of interdisciplinary dialogue in advancing both the scientific understanding and practical governance of multimodal AI.

We are deeply grateful to our workshop organisers, student organisers, program committee members, speakers, panelists, and authors for their time, expertise, and commitment to this effort. This workshop has been shaped by collaboration across institutions and countries, and we hope it offers a constructive space for rigorous discussion, new connections, and future research directions in responsible multimodal AI.

We warmly welcome you to the workshop and hope you find the program engaging, thought-provoking, and generative.

Soyeon Caren Han, Editors-in-Chief
Rina Carines Cabral, Editors-in-Chief

Josiah Poon, Editorial Board
Eduard Hovy, Editorial Board

Organising Committee

- **Soyeon Caren Han** – University of Melbourne, Australia
- **Rina Cabral** – University of Melbourne, Australia
- **Eduard Hovy** – University of Melbourne, Australia
- **Josiah Poon** – University of Sydney, Australia
- **Luca Cagliero** – Politecnico di Torino, Italy
- **Honghan Wu** – University of Glasgow, United Kingdom
- **Sang Wook Yi** – Hanyang University, South Korea
- **Kyungreem Han** – Korea Institute of Science and Technology, South Korea
- **Seungryong Kim** – Korea Advanced Institute of Science and Technology, South Korea
- **Kate Shim** – Korea Advanced Institute of Science and Technology, South Korea
- **Yihao Ding** – University of Western Australia, Australia
- **Minseok Song** – Pohang University of Science and Technology, South Korea
- **Goran Nenadic** – University of Manchester, United Kingdom
- **Christopher Leckie** – University of Melbourne, Australia
- **Prasenjit Mitra** – Carnegie Mellon University Africa, Rwanda

Student Organisers

- **David Chung** – University of Melbourne, Australia
- **Sehyuk Park** – University of Melbourne, Australia

Program Committee

- **Jie Yang** – China Mobile Research Institute, China
- **Sohyun Kang** – Korea Institute of Science and Technology, South Korea
- **Yerin Choi** – Korea Institute of Science and Technology, South Korea
- **Yue Dai** – University of Melbourne, Australia
- **Kalyana Krishna Kondapalli** – Mytecz
- **Dhita Putri Pratama** – University of Melbourne, Australia
- **Muku Akasaka** – University of Melbourne, Australia